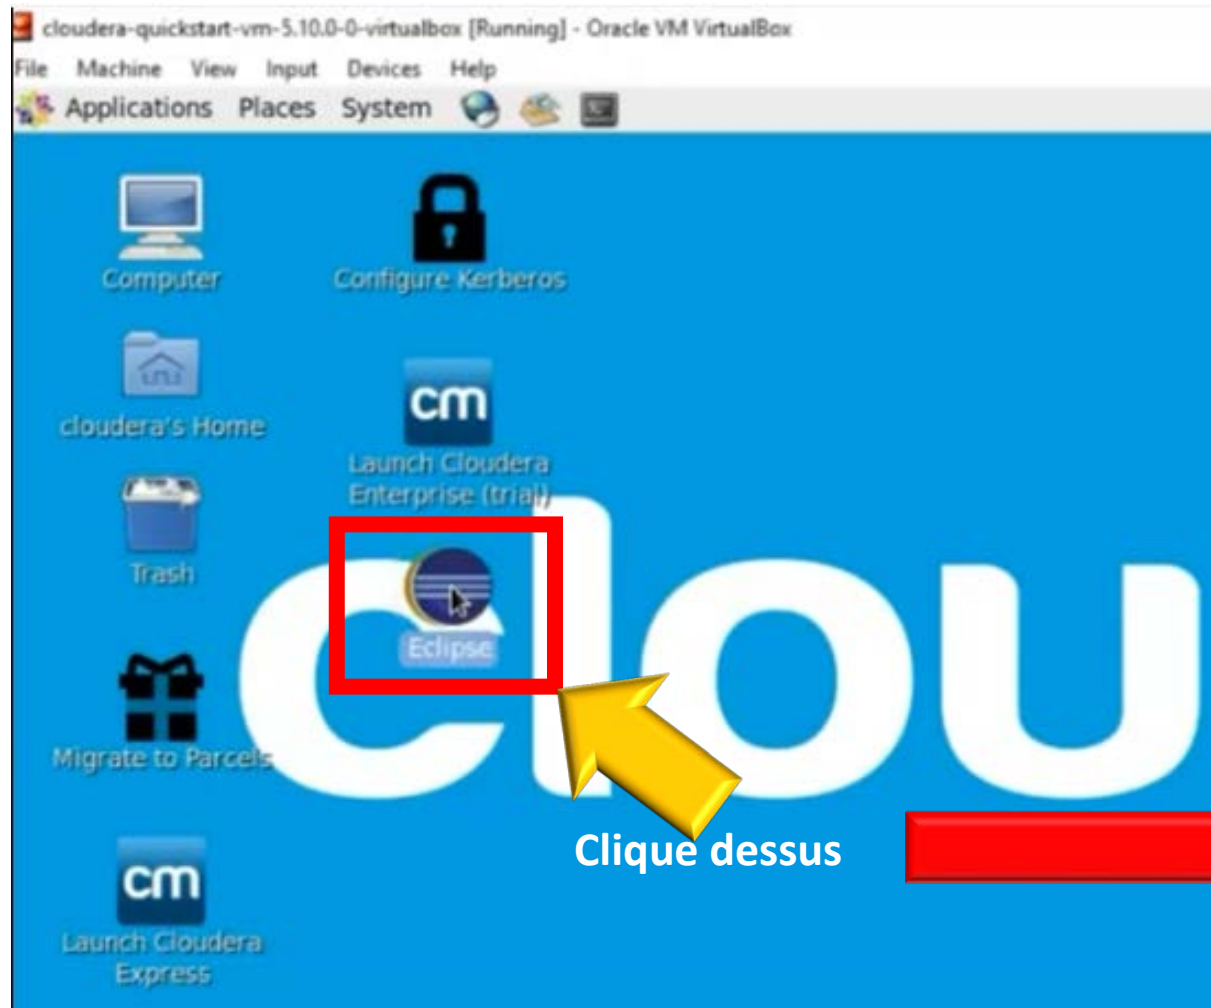


Atelier 1 Cloudera

Partie 2

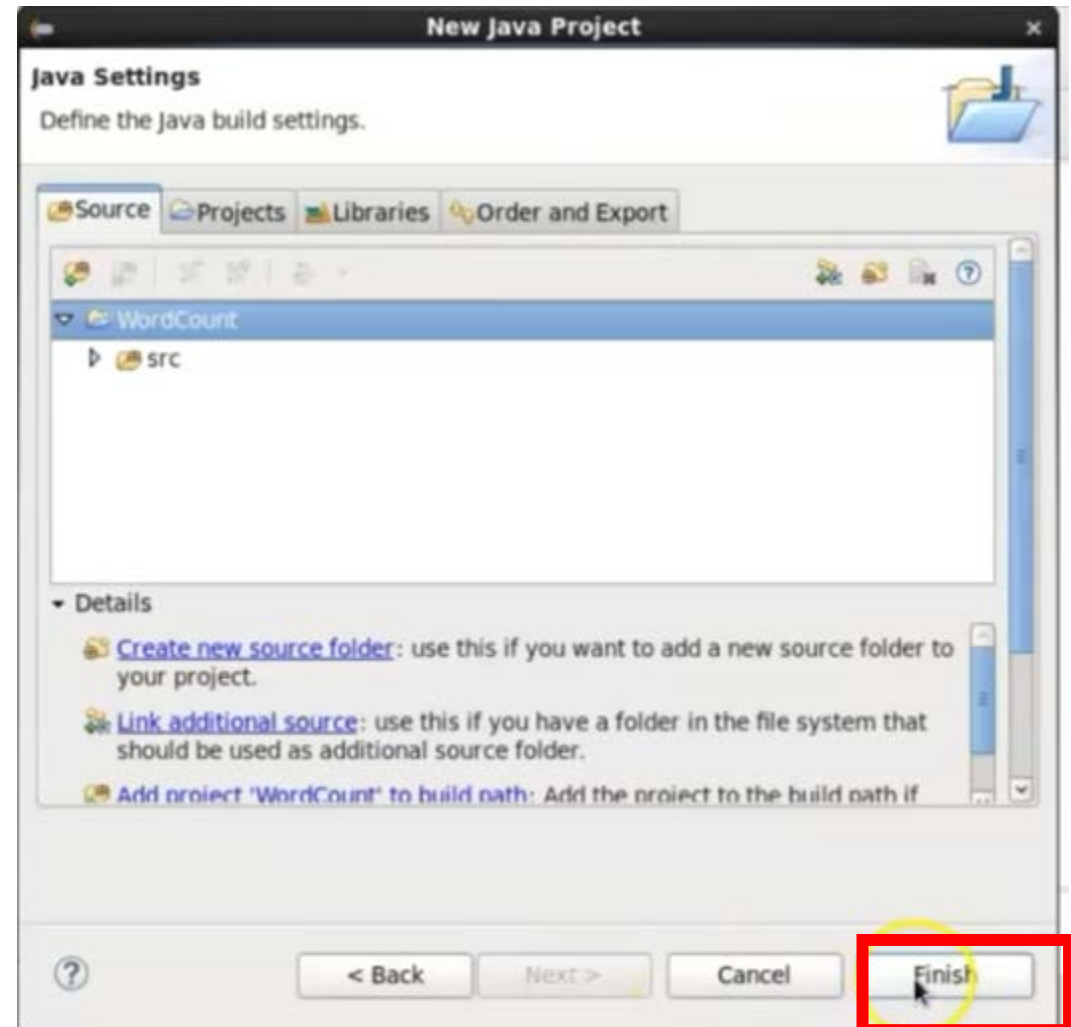
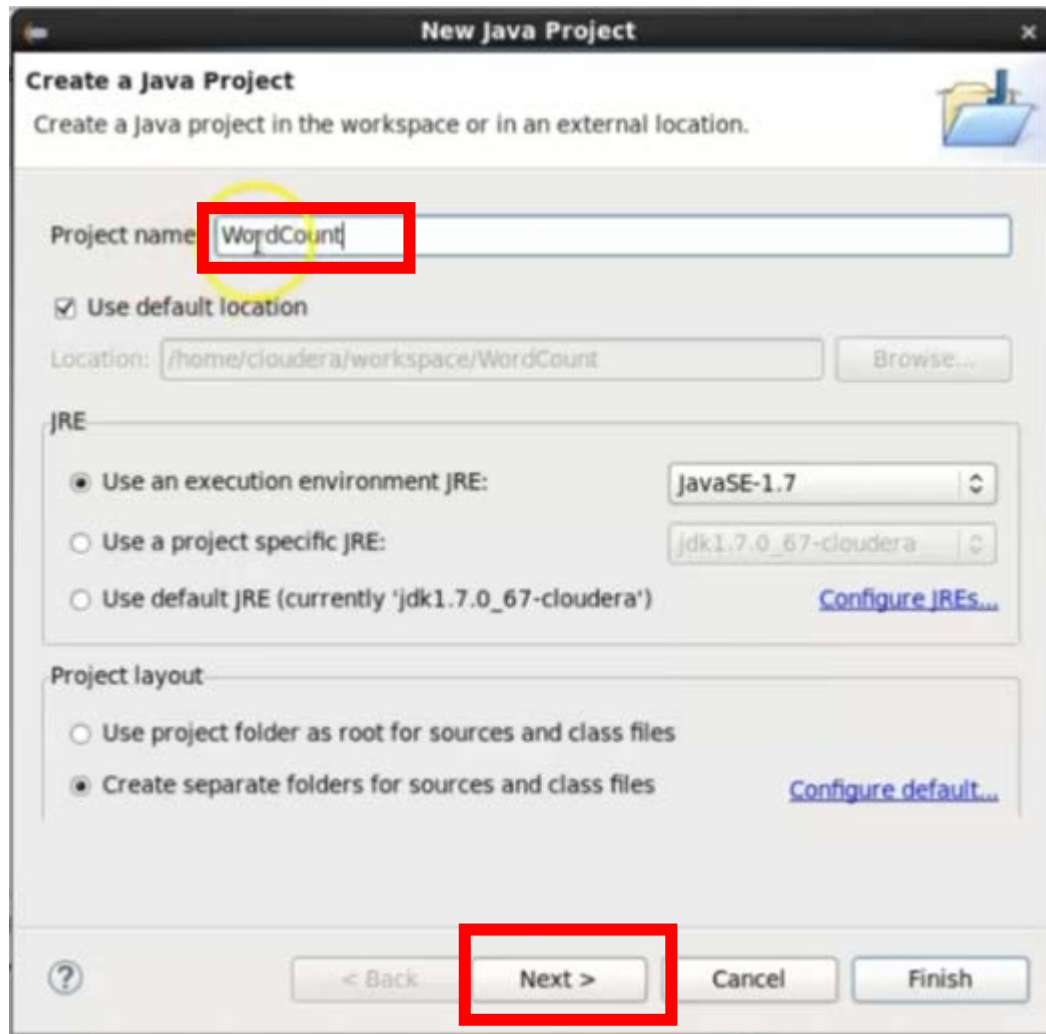
Map Reduce

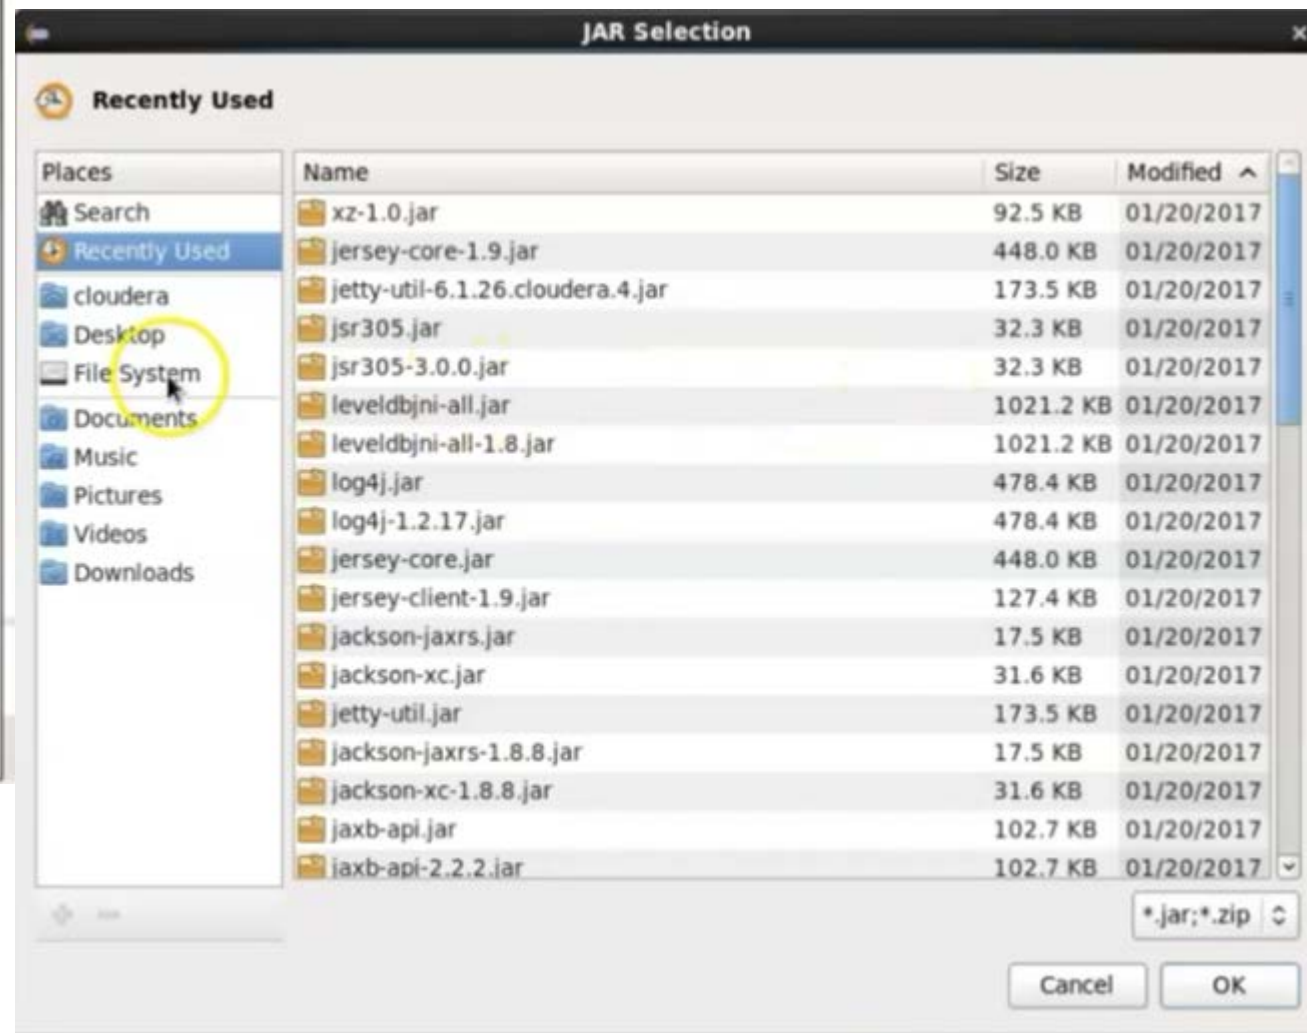



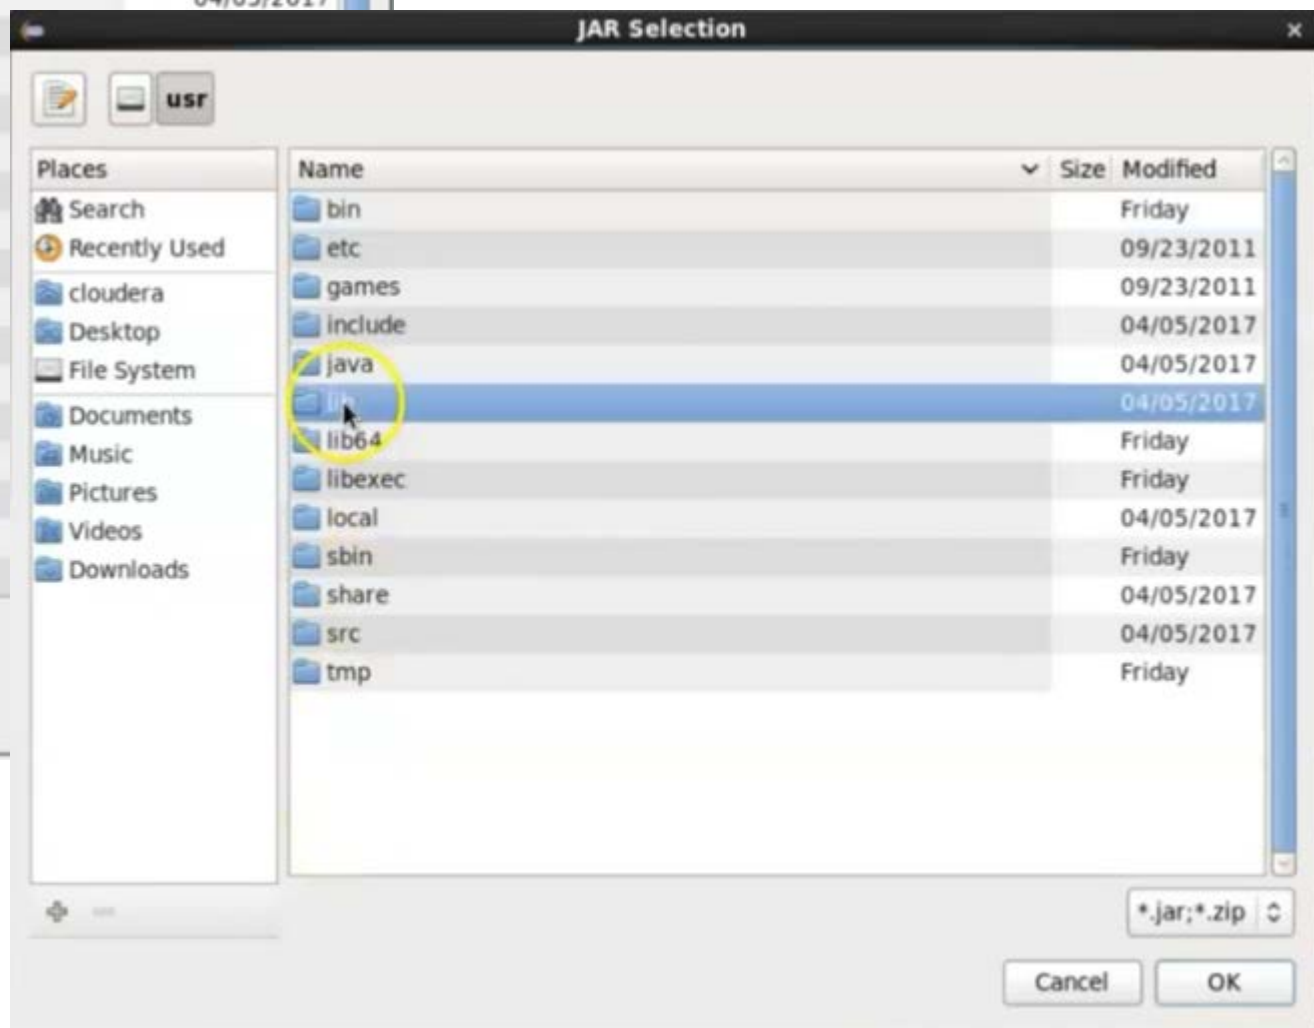
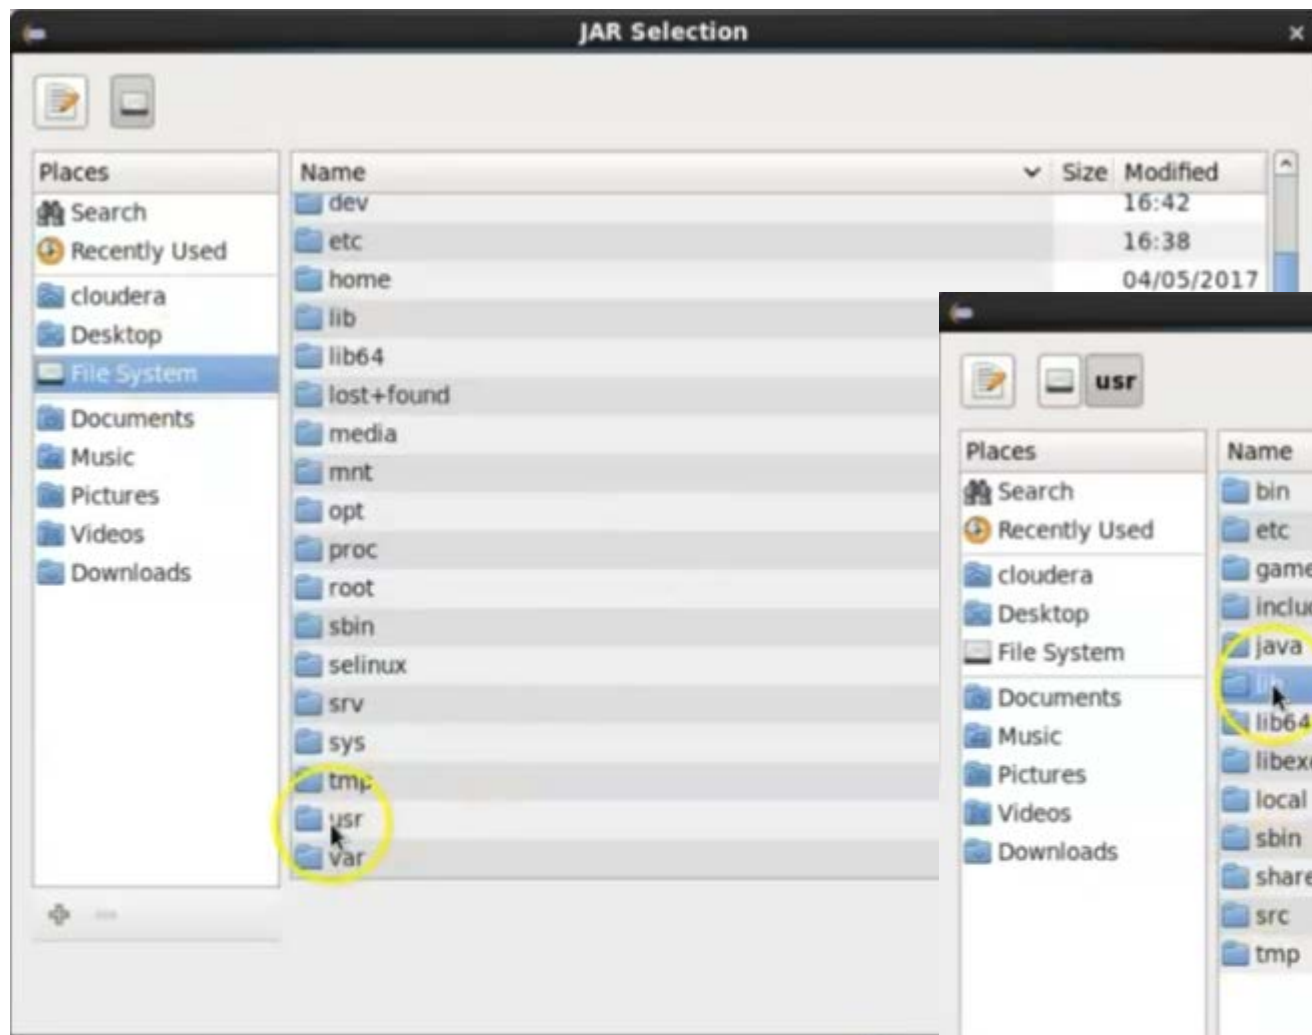
Java - Eclipse

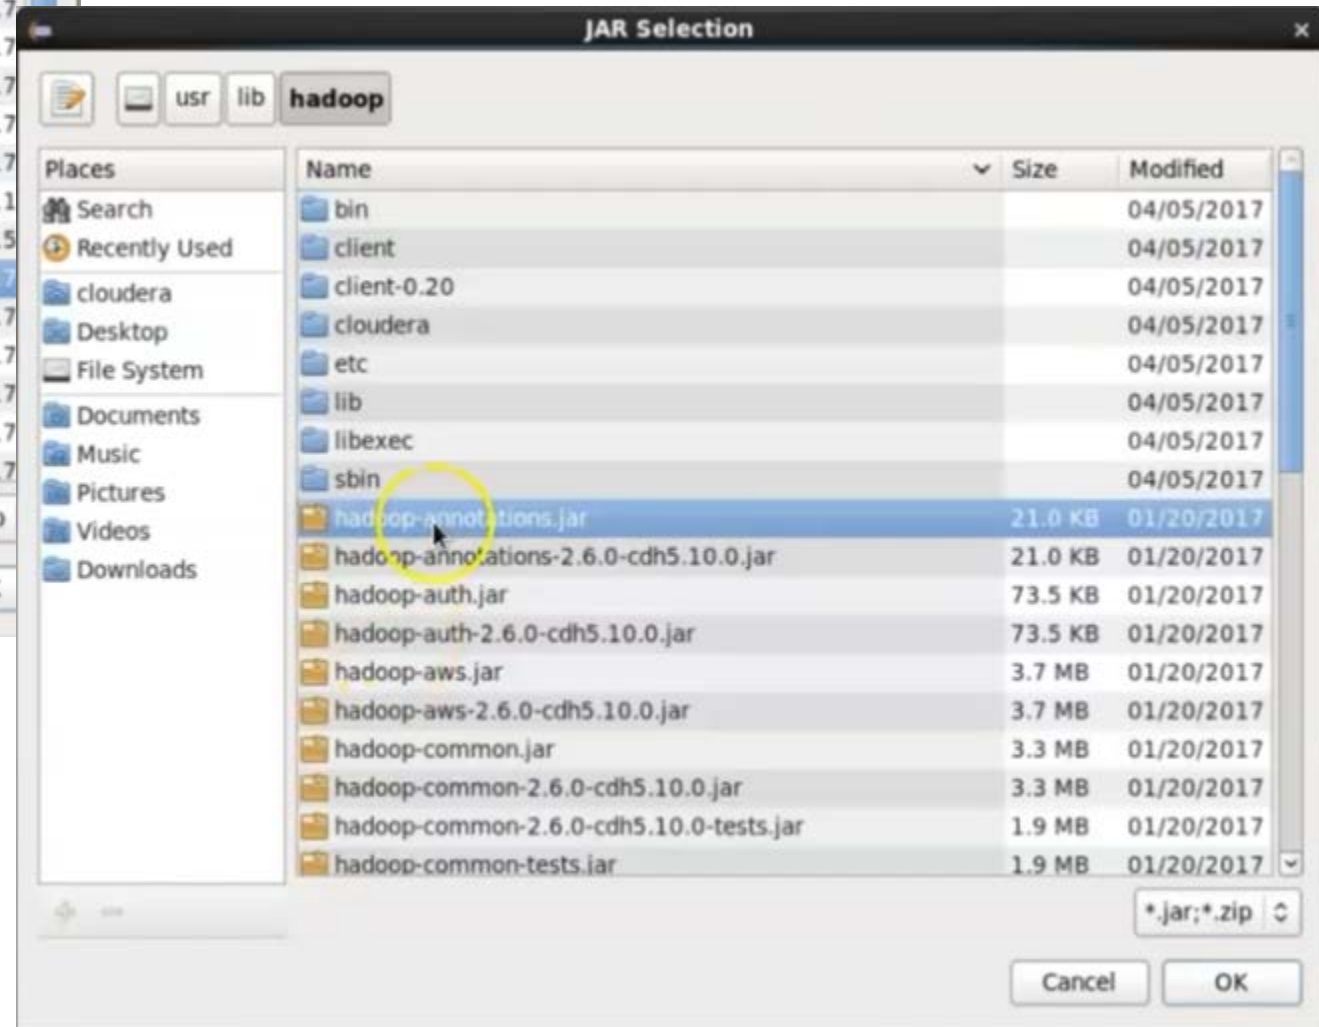
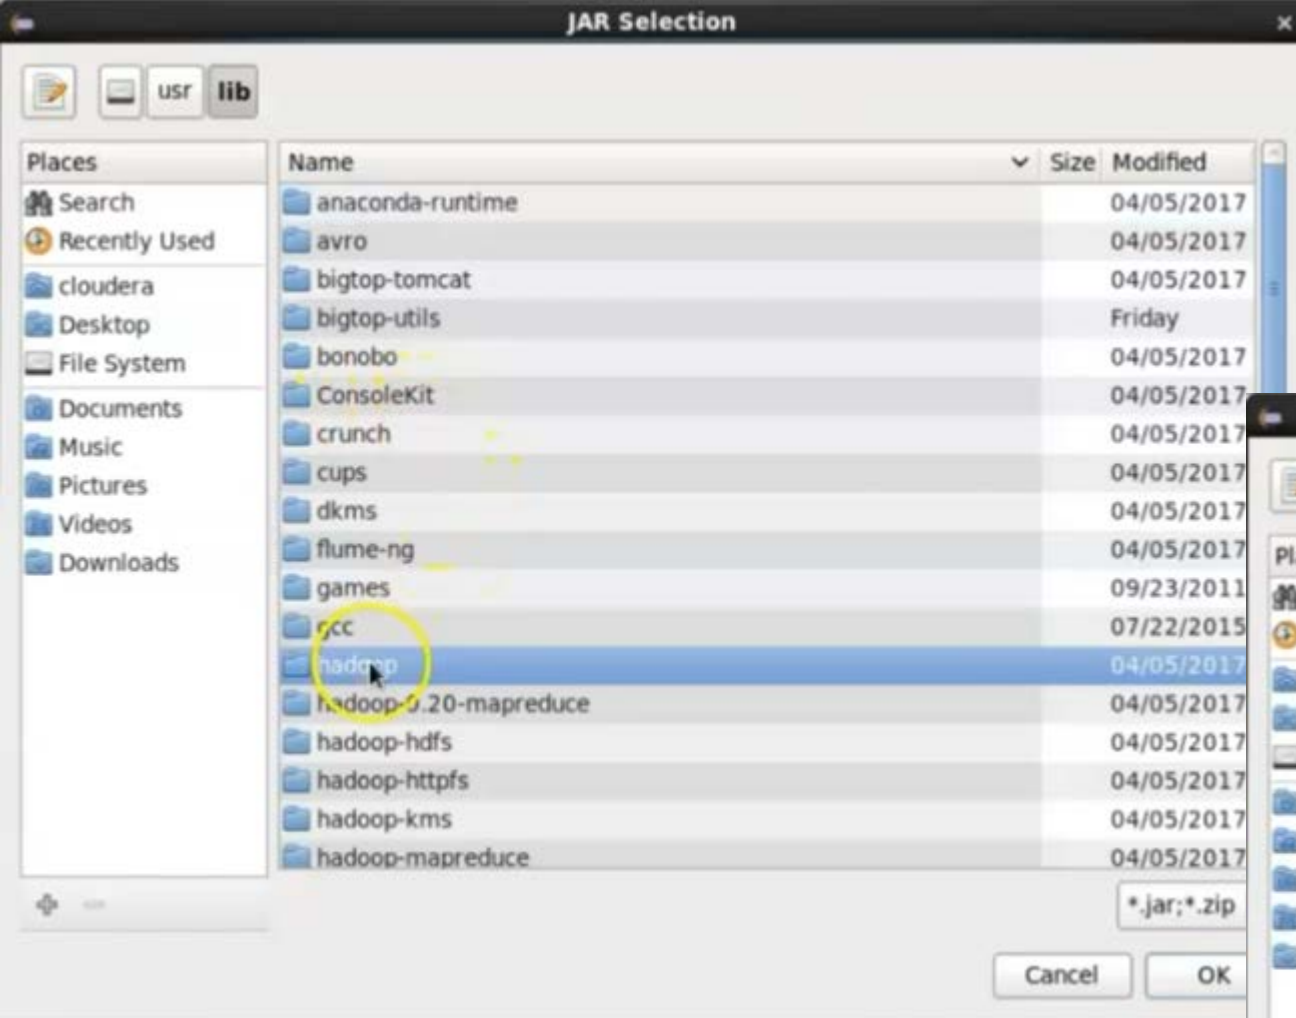
File Edit Source Refactor Navigate Search Project Run Window Help


New	Shift+Alt+N	Java Project
Open File...		Project...
Close	Ctrl+W	Package
Close All	Shift+Ctrl+W	Class
Save	Ctrl+S	Interface
Save As...		Enum
Save All	Shift+Ctrl+S	Annotation
Revert		Source Folder
Move...		Java Working Set
Rename...	F2	Folder
Refresh	F5	File
Convert Line Delimiters To	>	Untitled Text File
Print...	Ctrl+P	JUnit Test Case
Switch Workspace	>	Task
Restart		Example...
Import...		Other... Ctrl+N
Export...		





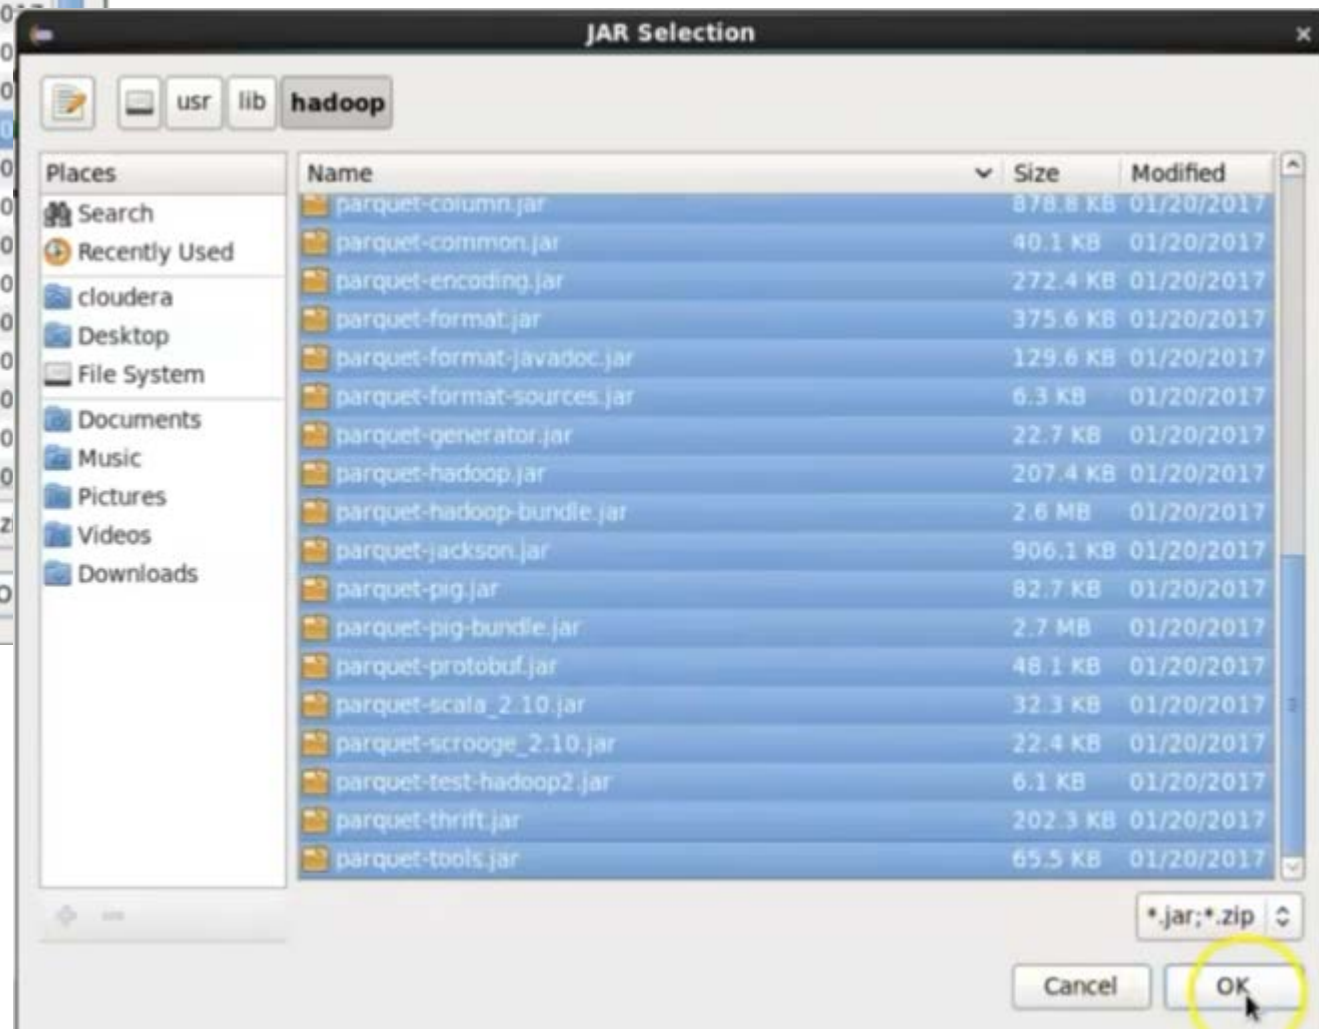



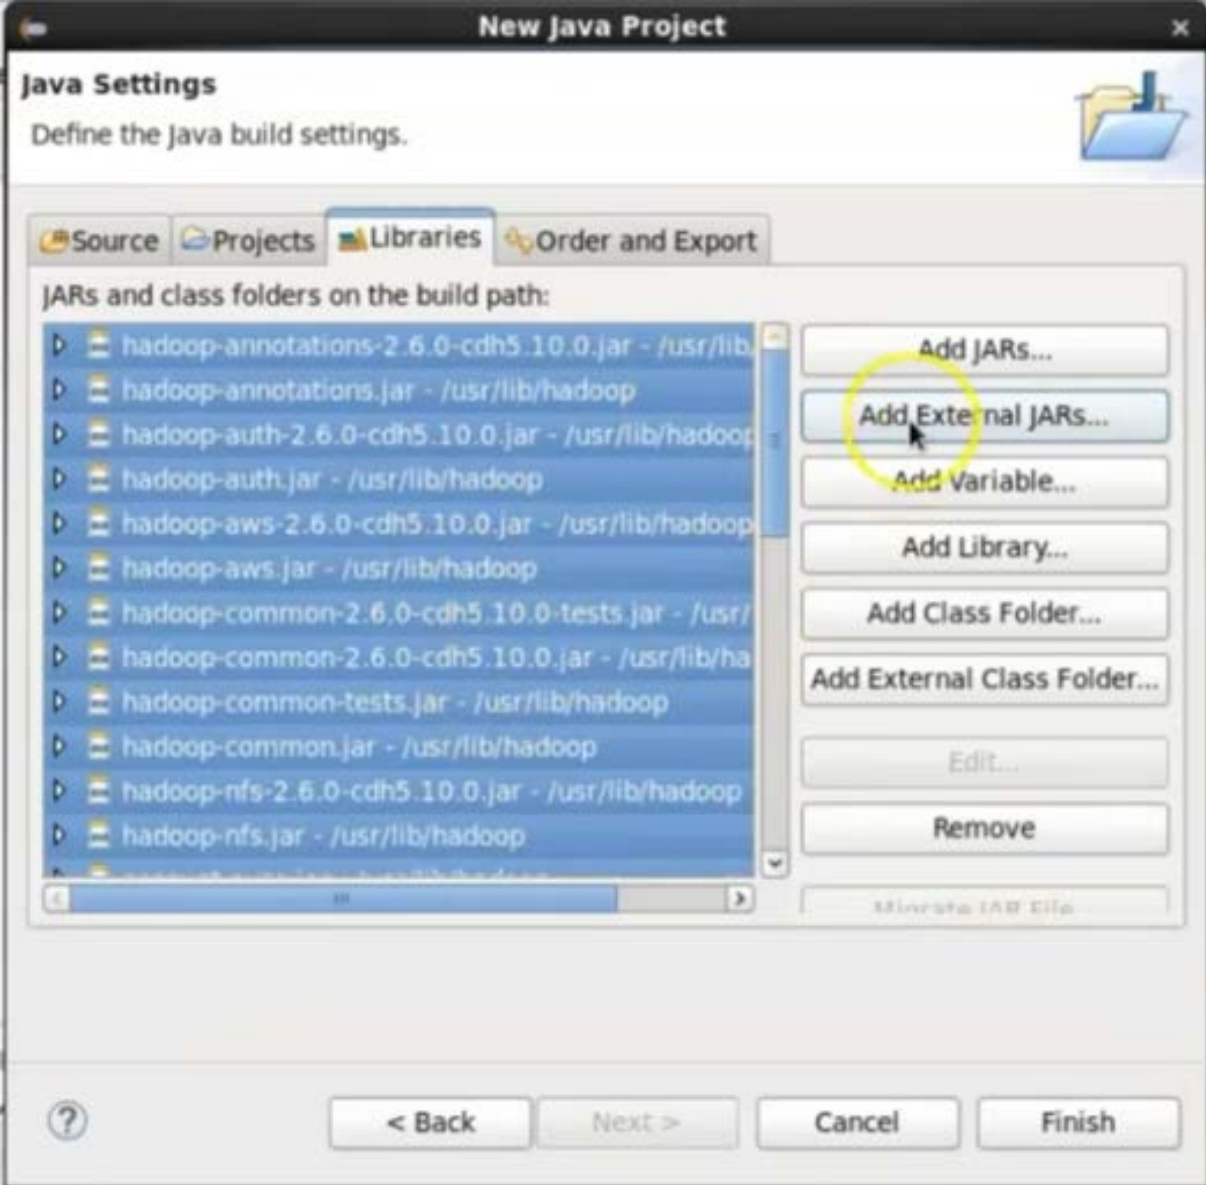


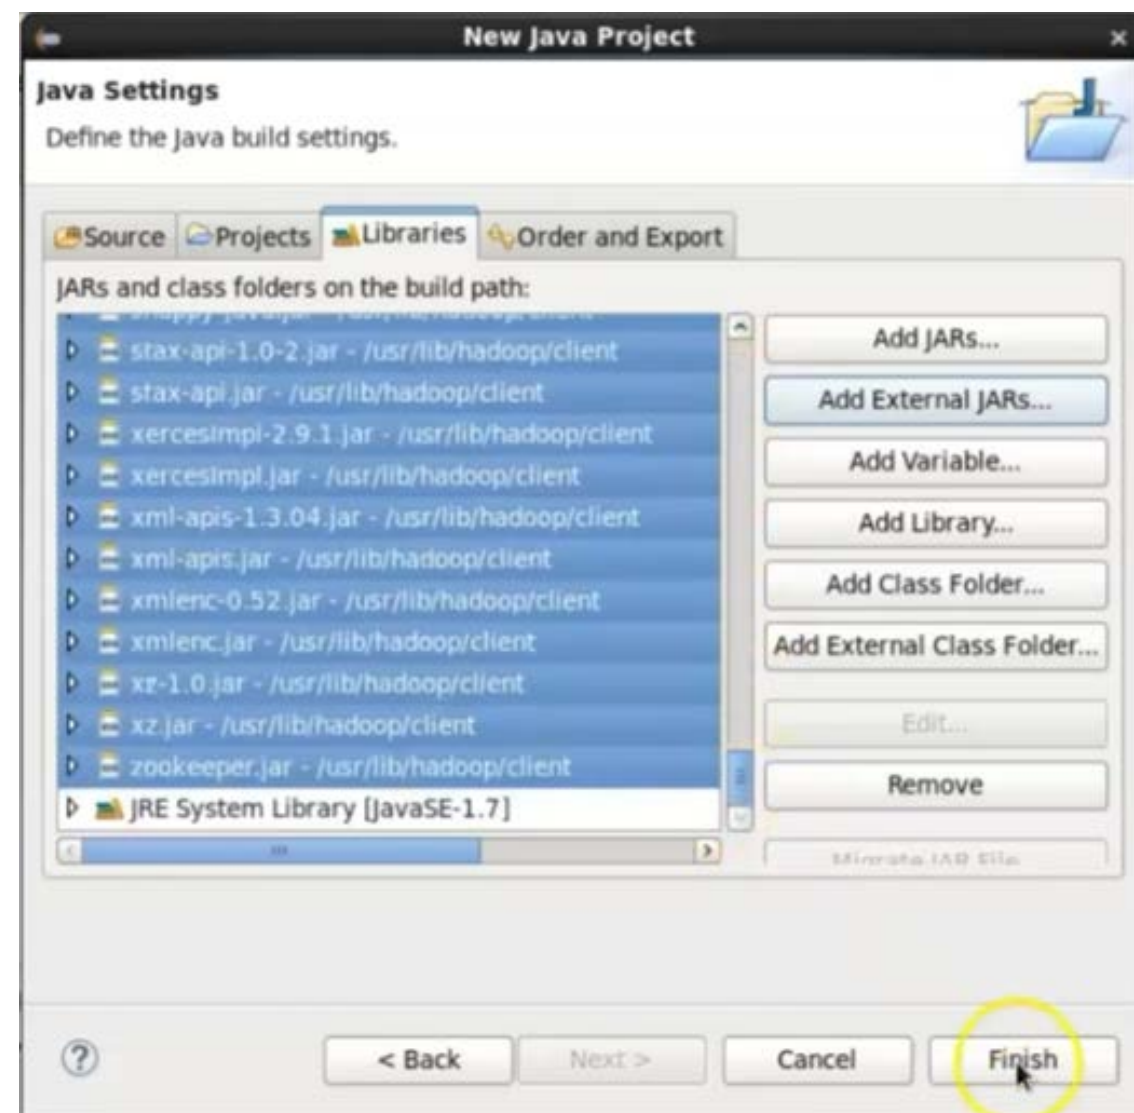
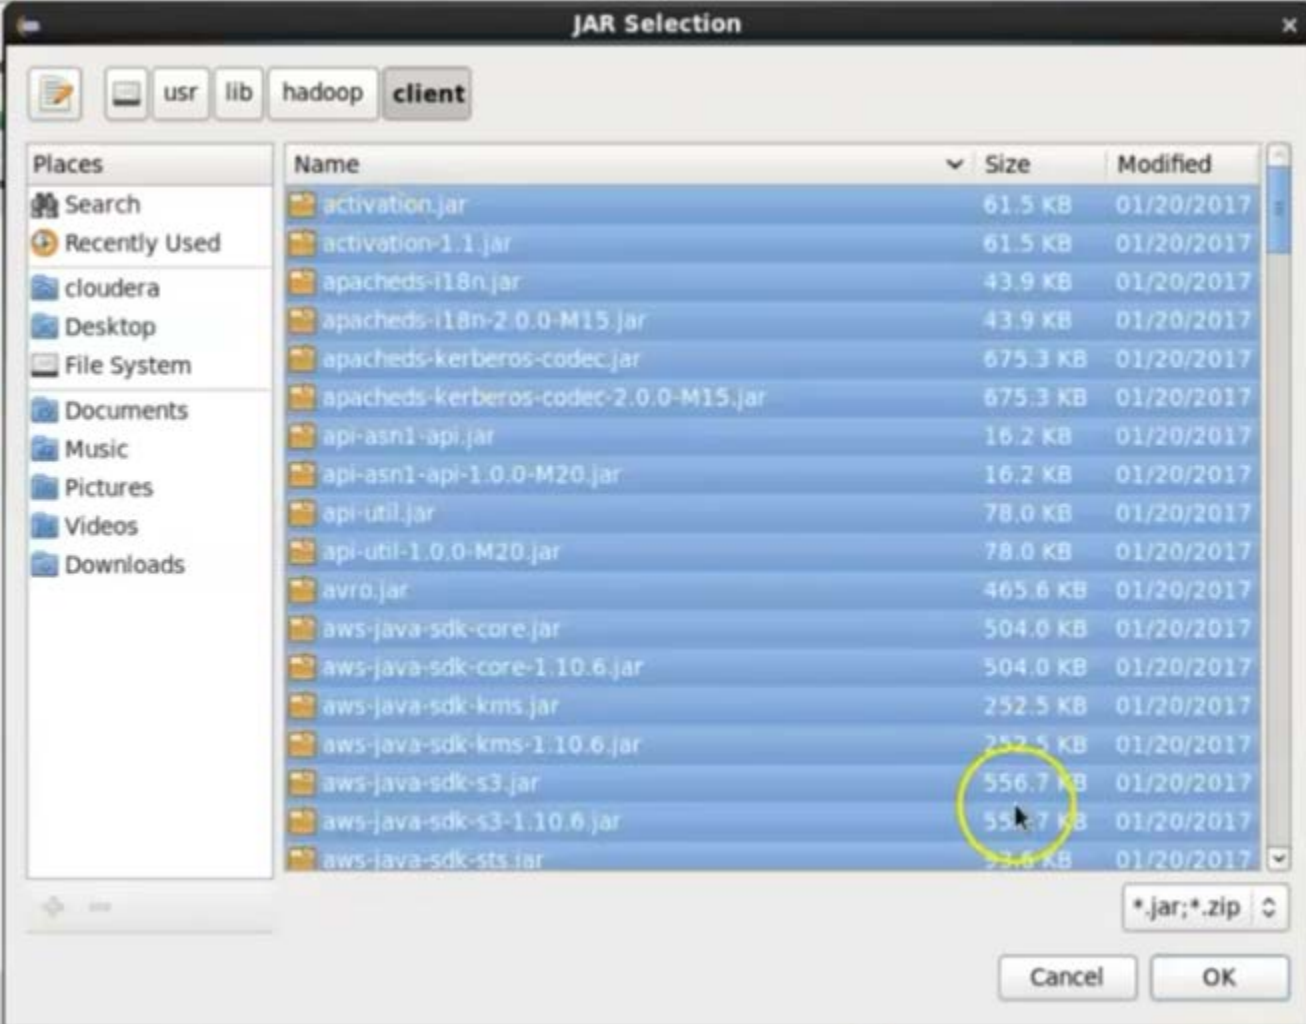


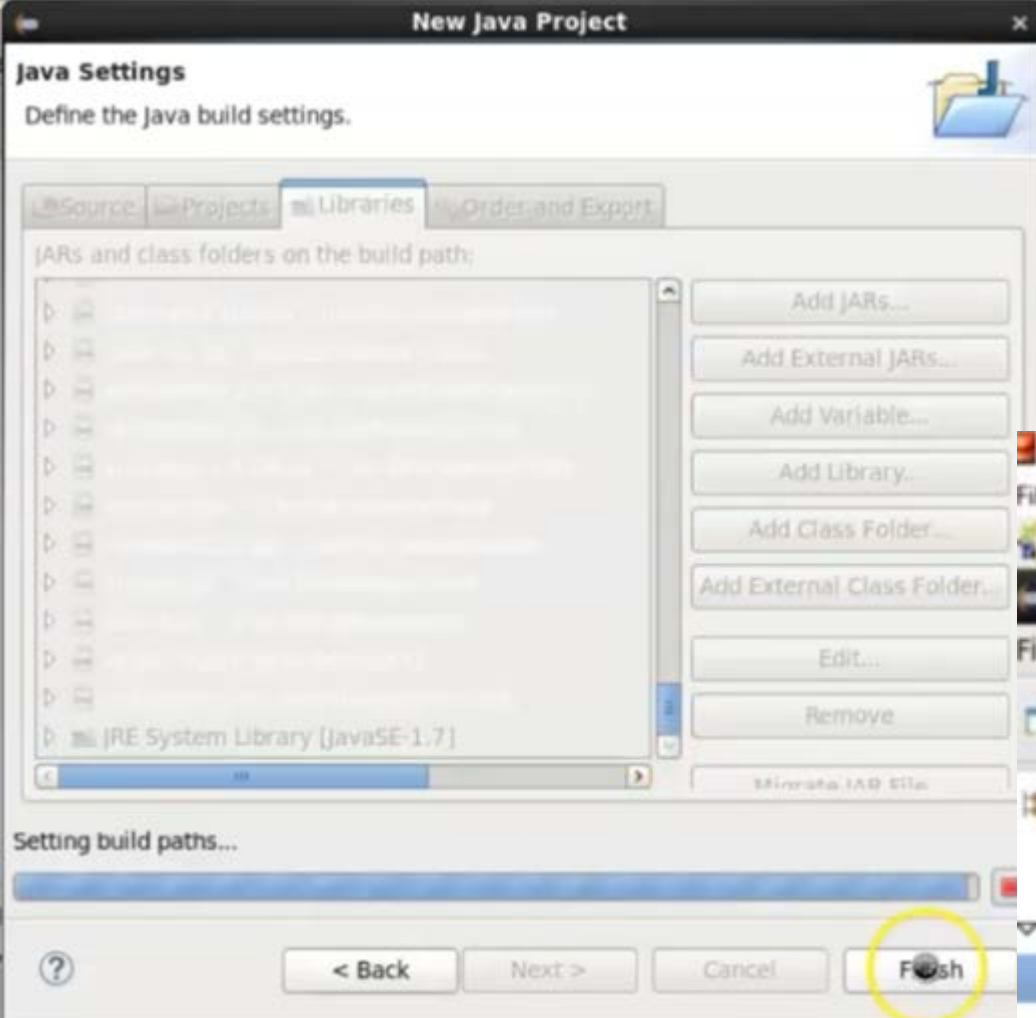
NB :

Sélection la colonne de : hadoop annotation jusqu'a le bas parquet-tools.jar









cloudera-quickstart-vm-5.10.0-0-virtualbox [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Applications Places System

Java - Eclipse

File Edit Source Refactor Navigate Search Project Run Window Help

Package Explorer

WordCount

JRE System Library [javaSE-1.7]

Referenced Libraries

Java - Eclipse

File Edit New Source Refactor Navigate Search Project Run Window Help

New

- Open in New Window
- Open Type Hierarchy F4
- Show In Shift+Alt+W
- Copy Ctrl+C
- Copy Qualified Name
- Paste Ctrl+V
- Delete Delete
- Remove from Context Shift+Ctrl+Alt+Down
- Build Path
- Source Shift+Alt+S
- Refactor Shift+Alt+T
- Import...
- Export...
- Refresh F5
- Assign Working Sets...
- Debug As
- Run As
- Validate
- Team
- Compare With
- Restore from Local History...
- Maven
- Properties Alt+Enter

Java Project
Project...

Package

Class

Interface

Enum

Annotation

Source Folder

Java Working Set

Folder

File

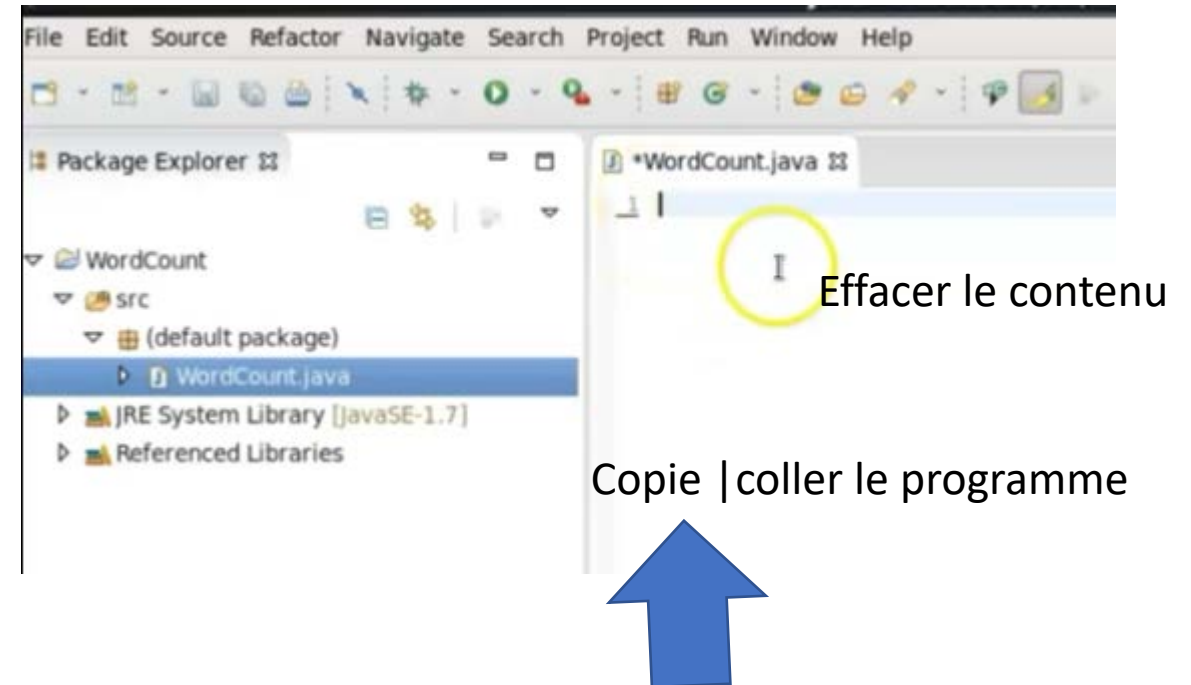
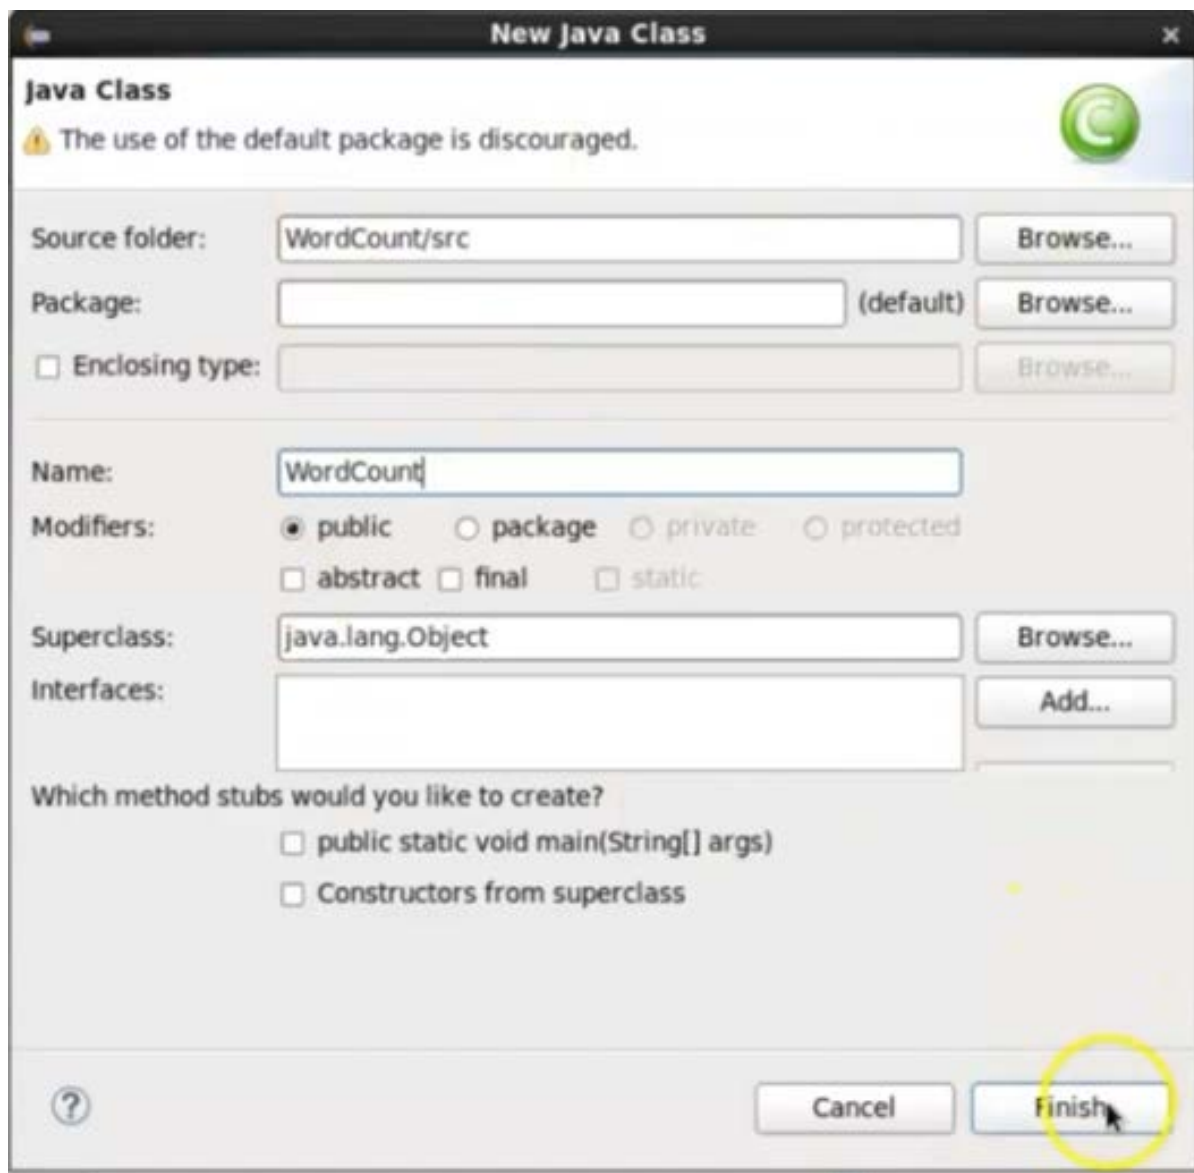
Untitled Text File

JUnit Test Case

Task

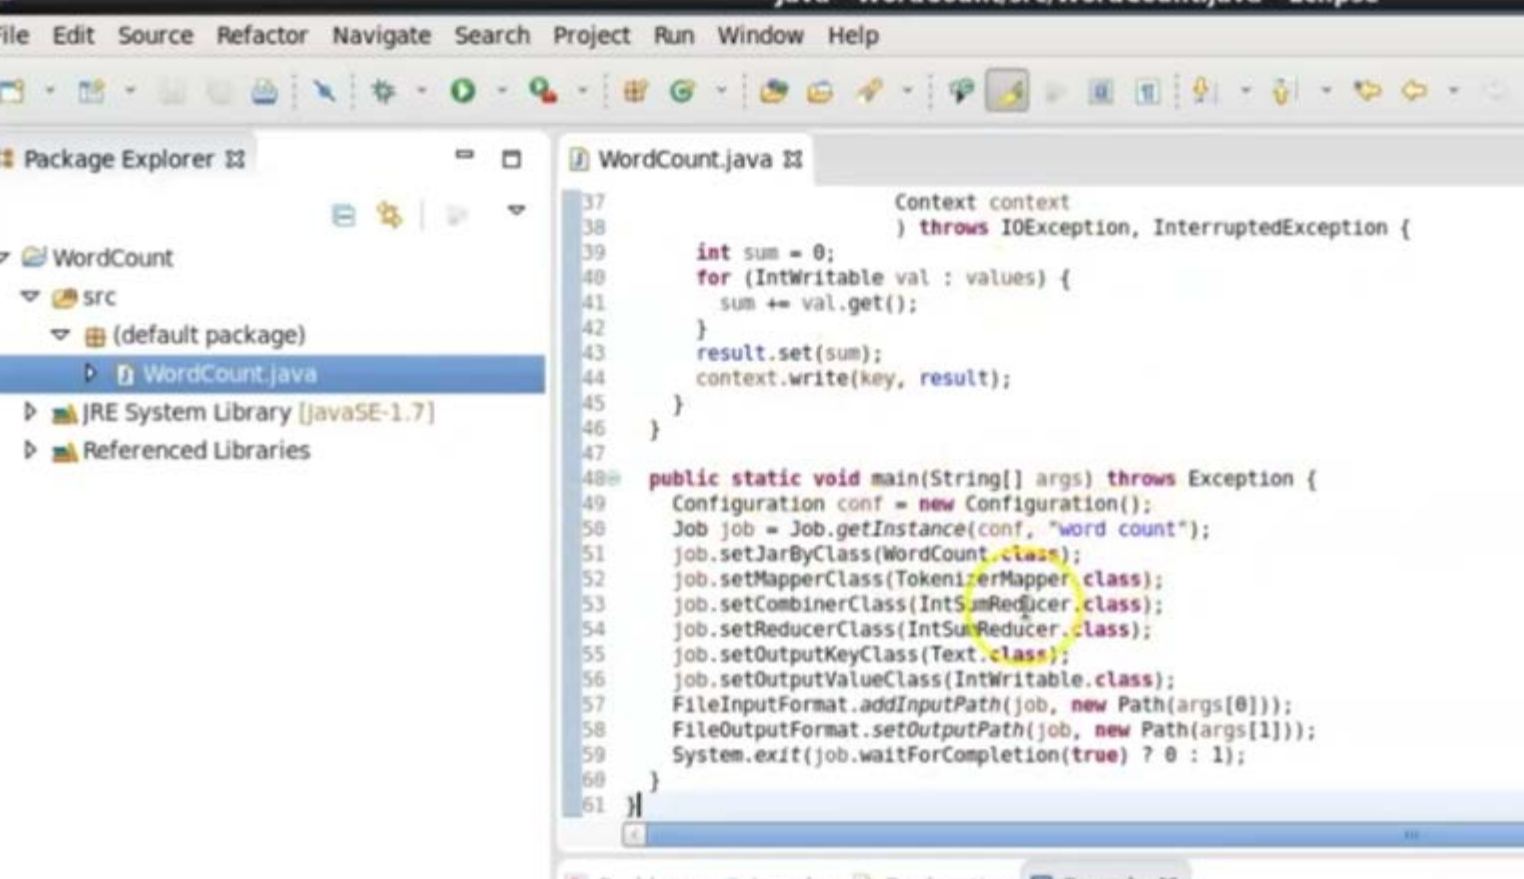
Example...

Other... Ctrl+N



<https://hadoop.apache.org/docs/stable/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>

Voir référence



e VM VirtualBox

Mon May 4, 5:06 PM cloudera

Apache Hadoop 3.2.1 - MapReduce Tutorial - Mozilla Firefox

oop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceT

Search

Impala Spark Solr Oozie Cloudera Manager Getting Started

- Building
- Tools
 - Hadoop Streaming
 - Hadoop Archives
 - Hadoop Archive Logs
 - DistCp
 - GridMix
 - Runen
 - Resource Estimator Service
 - Scheduler Load Simulator
 - Hadoop Benchmarking
- Reference
 - Changelog and Release Notes
 - Java API docs
 - Unix Shell API
 - Metrics
- Configuration
 - core-default.xml
 - hdfs-default.xml
 - hdfs-rbf-default.xml
 - mapred-default.xml
 - yam-default.xml
 - Deprecated Properties

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
```

Mozilla Firefox seems slow... to... start.

Learn How to Sp

New
Go Into
Open in New Window
Open Type Hierarchy F4
Show In Shift+Alt+W
Copy Ctrl+C
Copy Qualified Name
Paste Ctrl+V
Delete Delete
Remove from Context Shift+Ctrl+Alt+Down
Build Path
Source Shift+Alt+S
Refactor Shift+Alt+T
Import...
Export...
Refresh F5
Close Project
Assign Working Sets...
Debug As
Run As
Validate
Team
Compare With
Restore from Local History...
Configure
Properties Alt+Enter

Java - WordCount/src/WordCount.java - Eclipse

Window Help

WordCount.java

```
static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

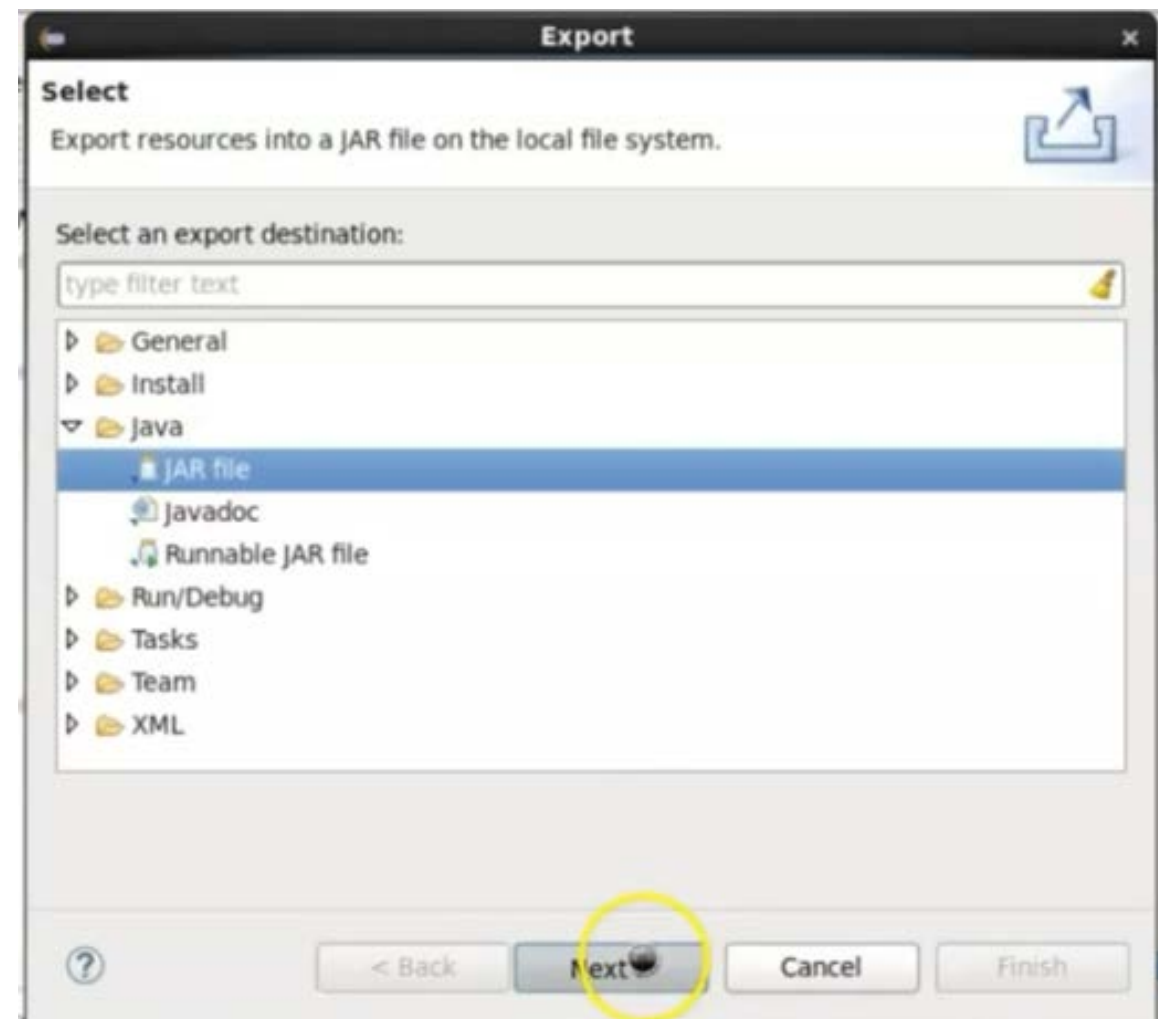
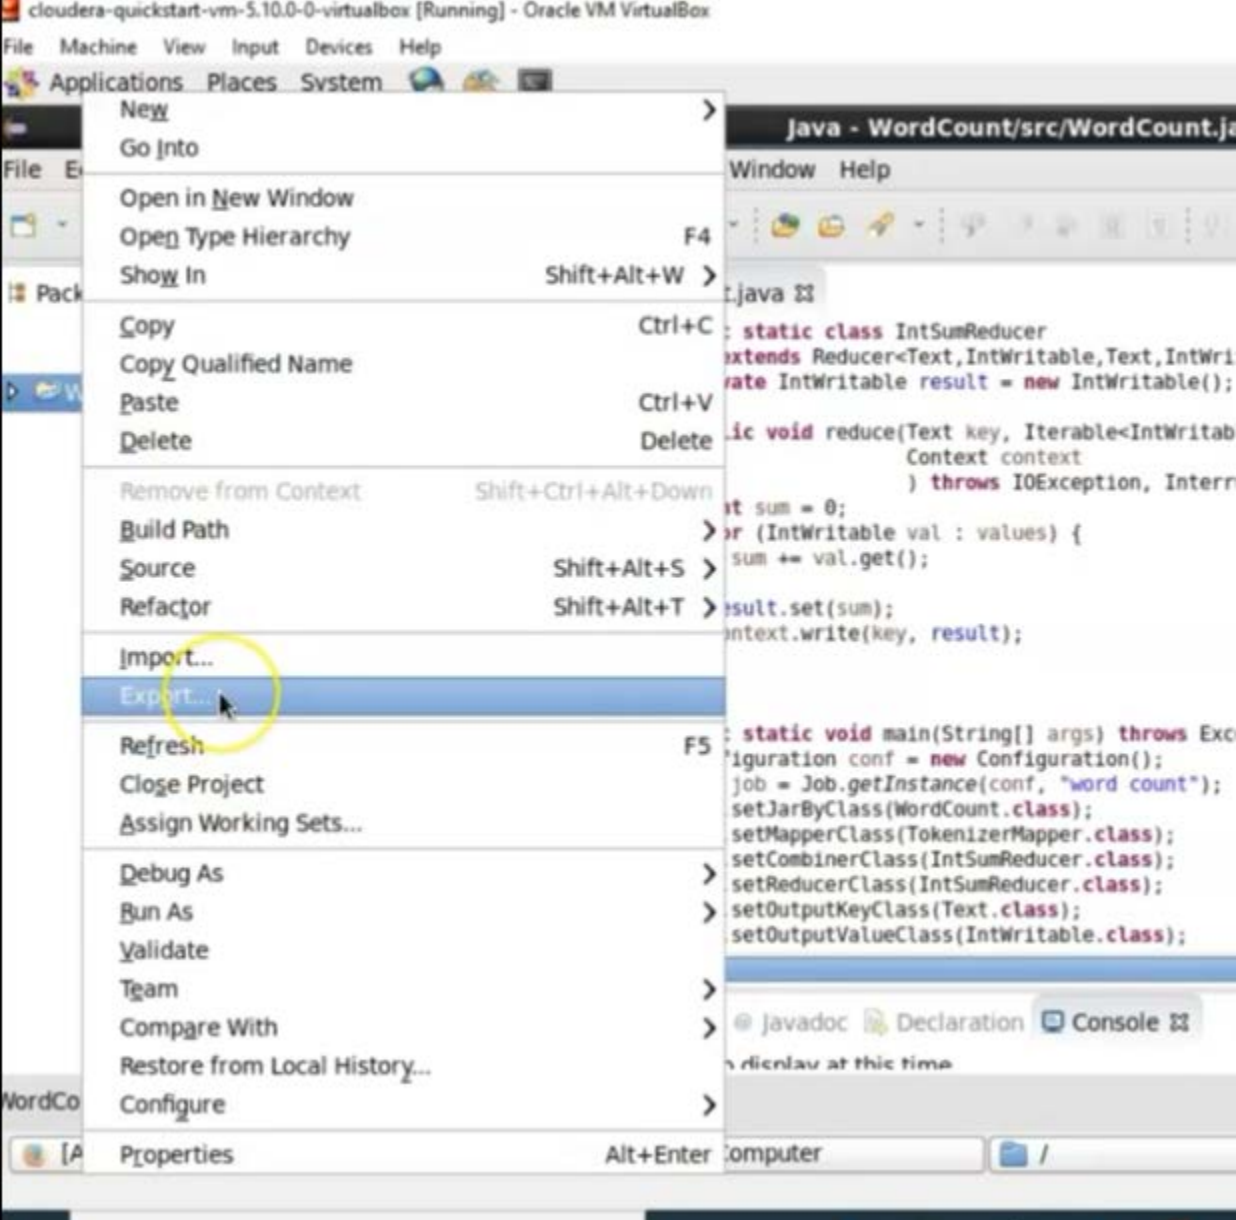
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
    ) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```


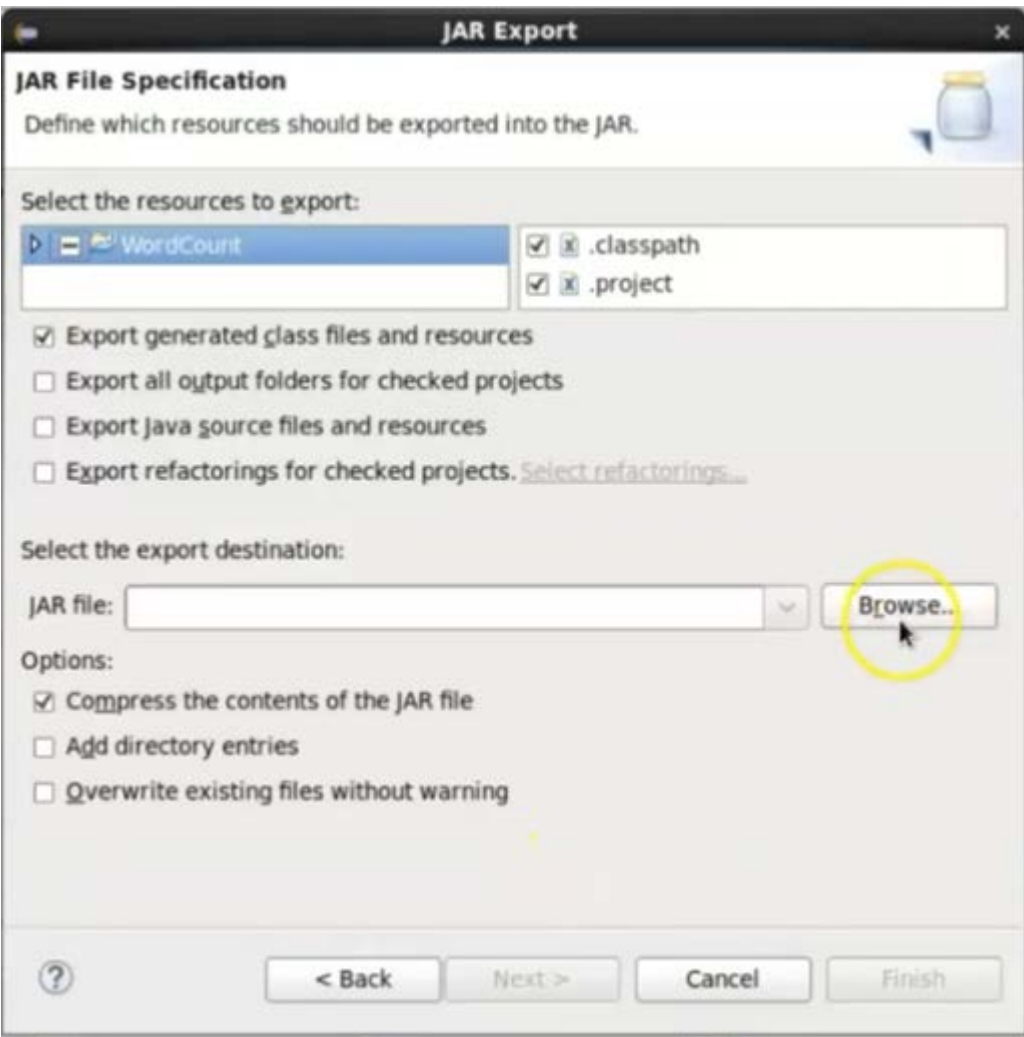
```
static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
}
```

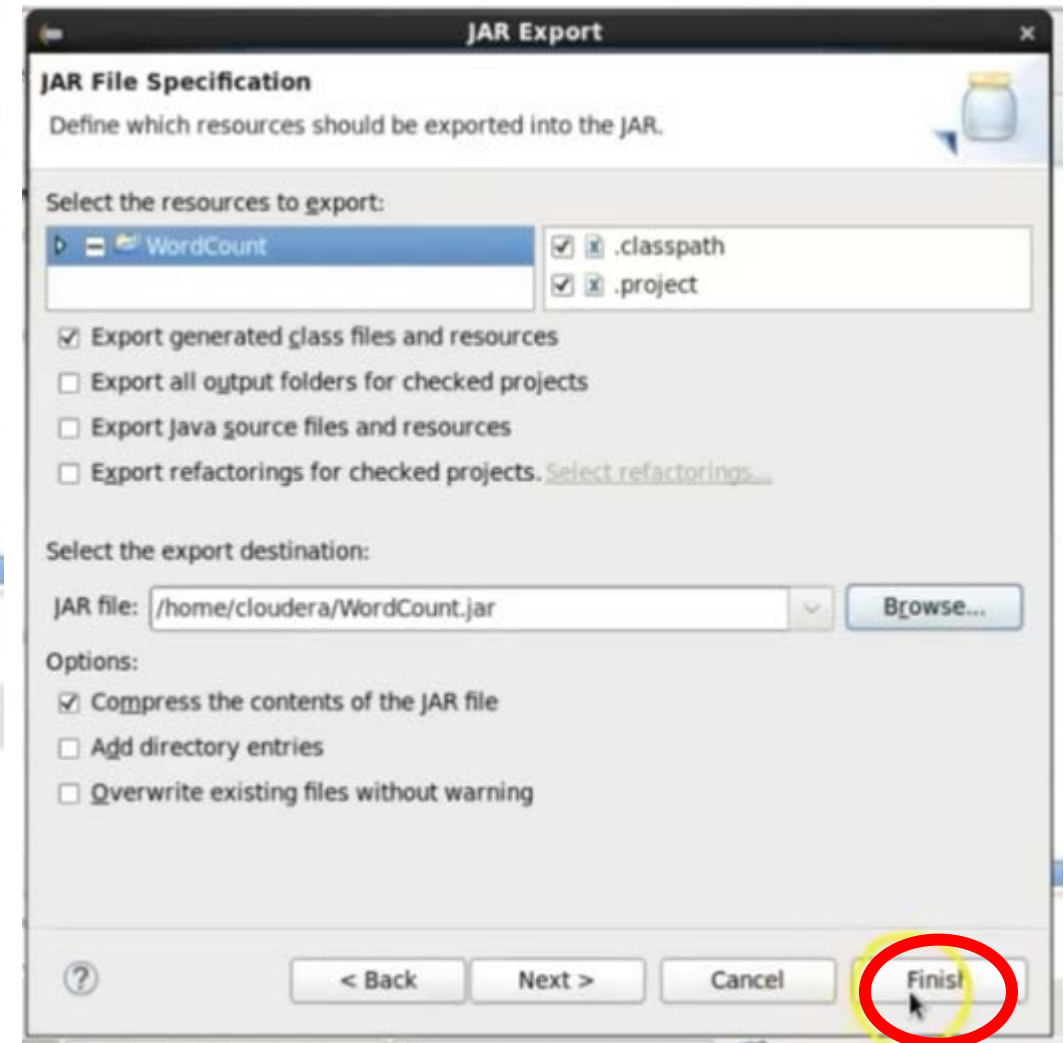
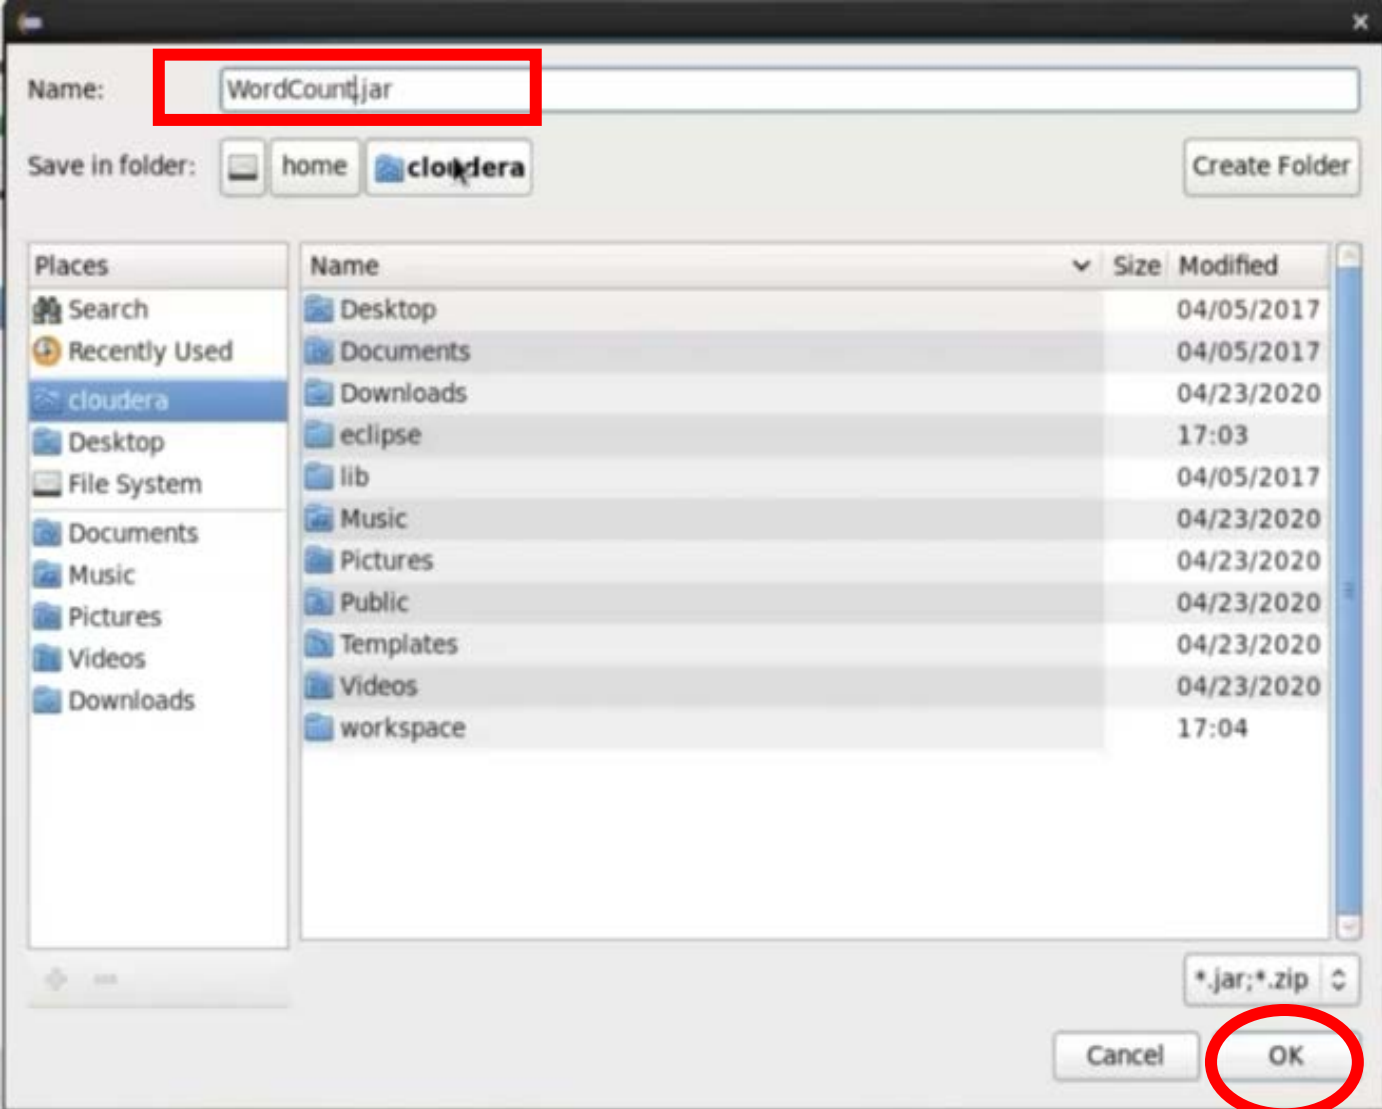
Javadoc Declaration Console

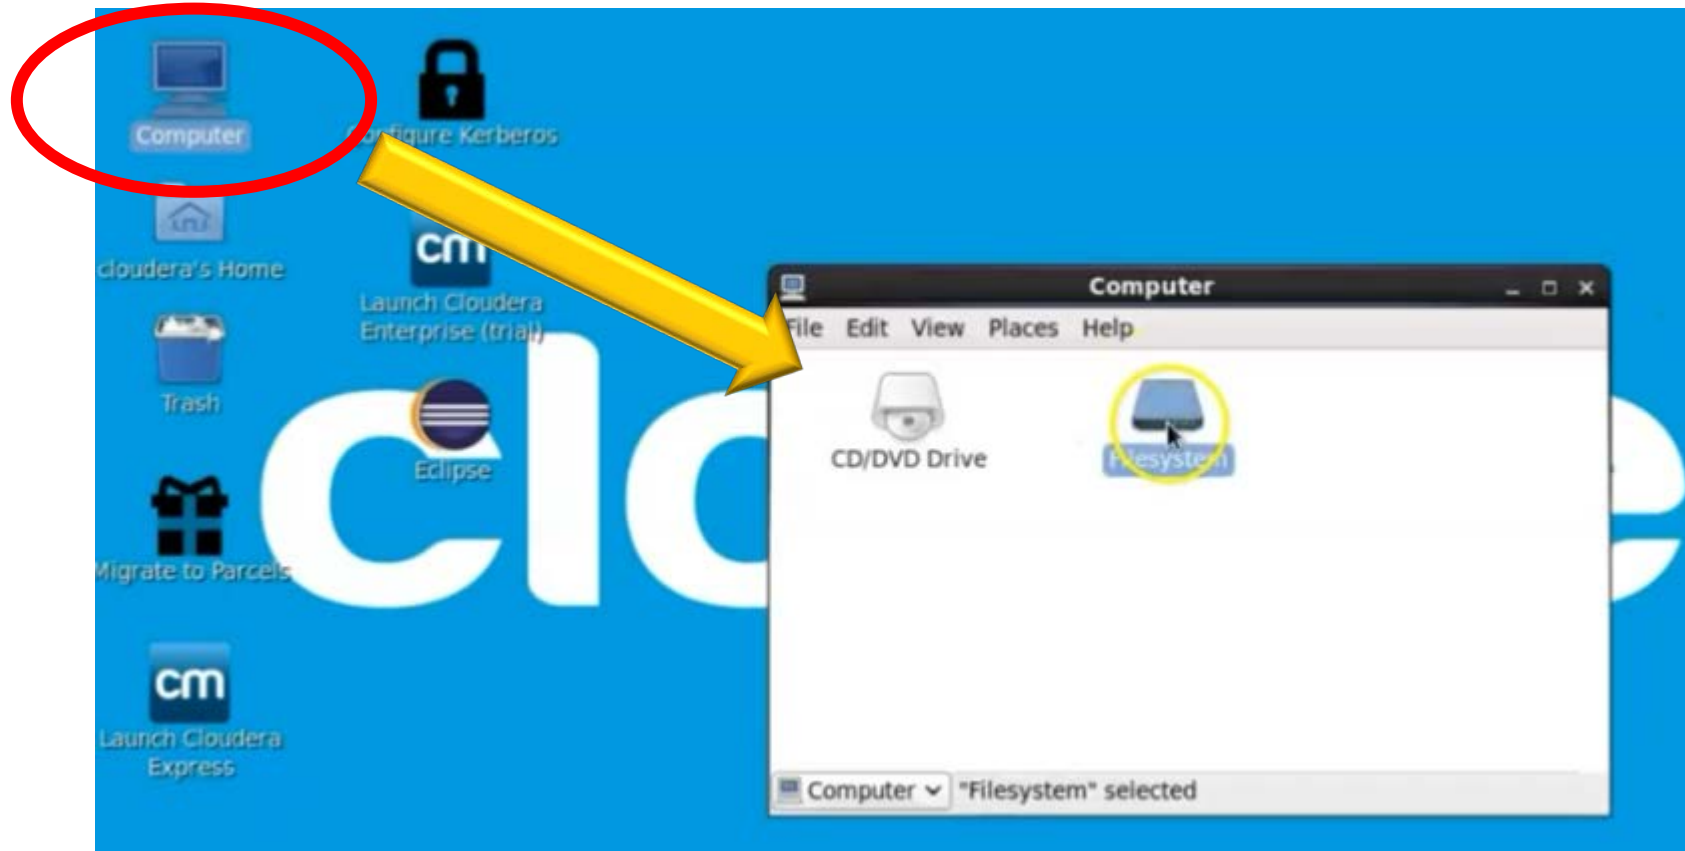
WordCount at this time

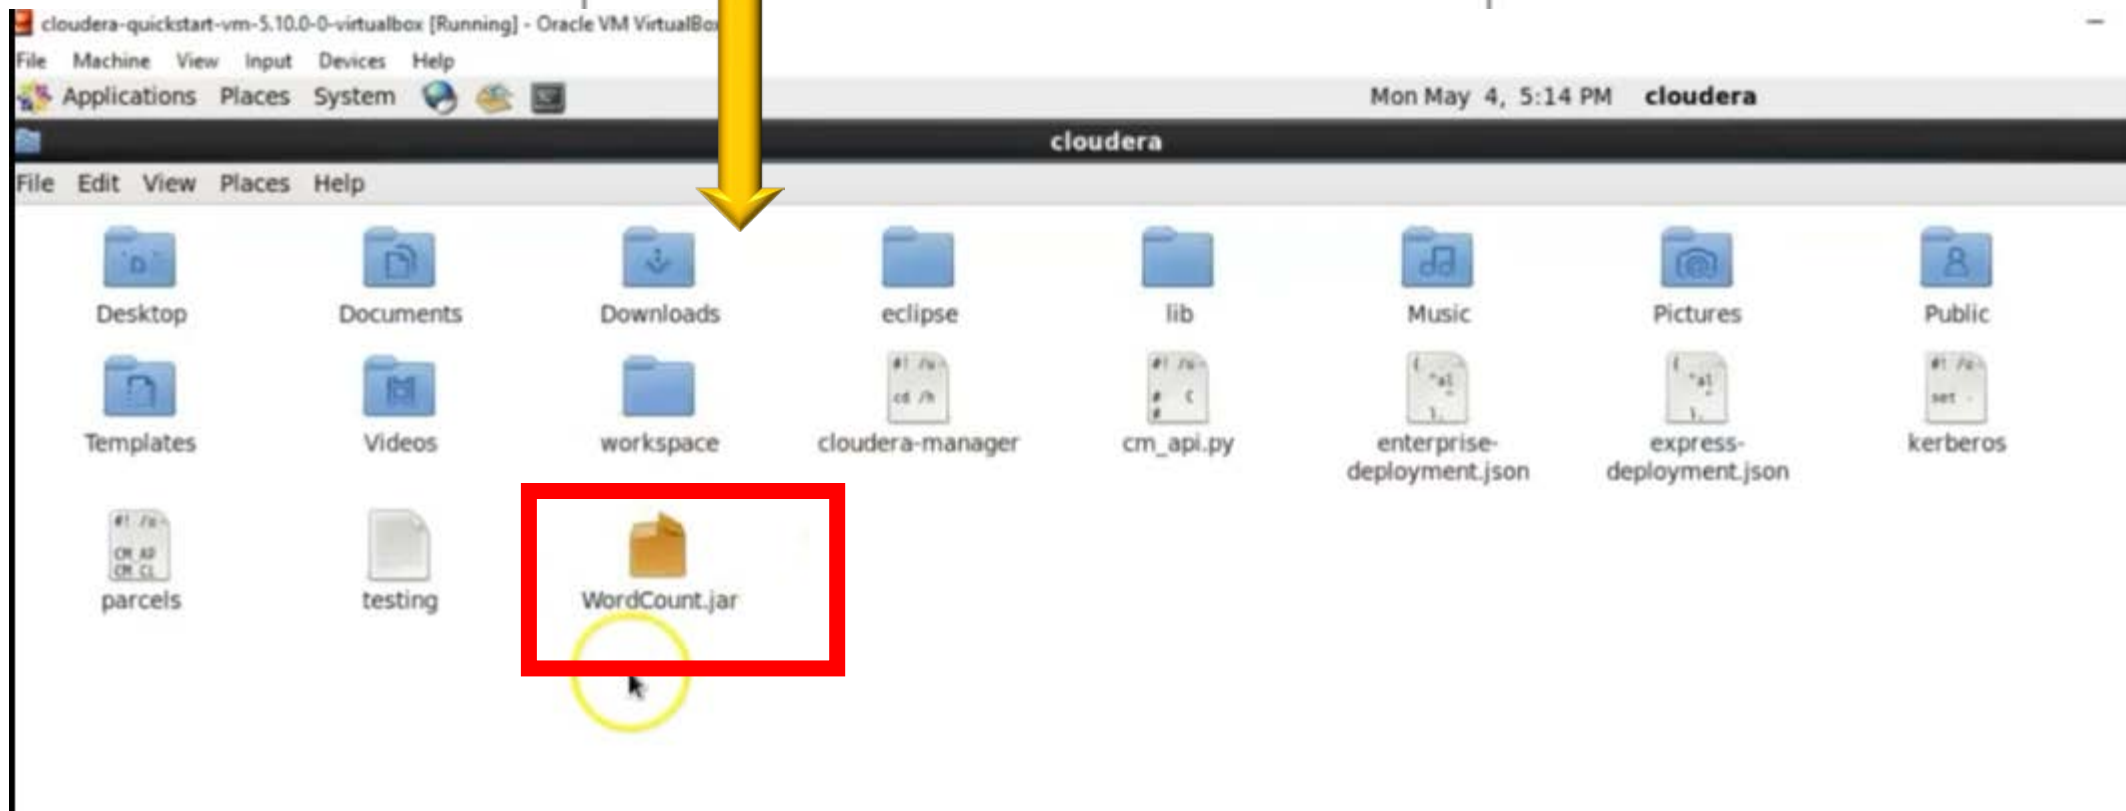
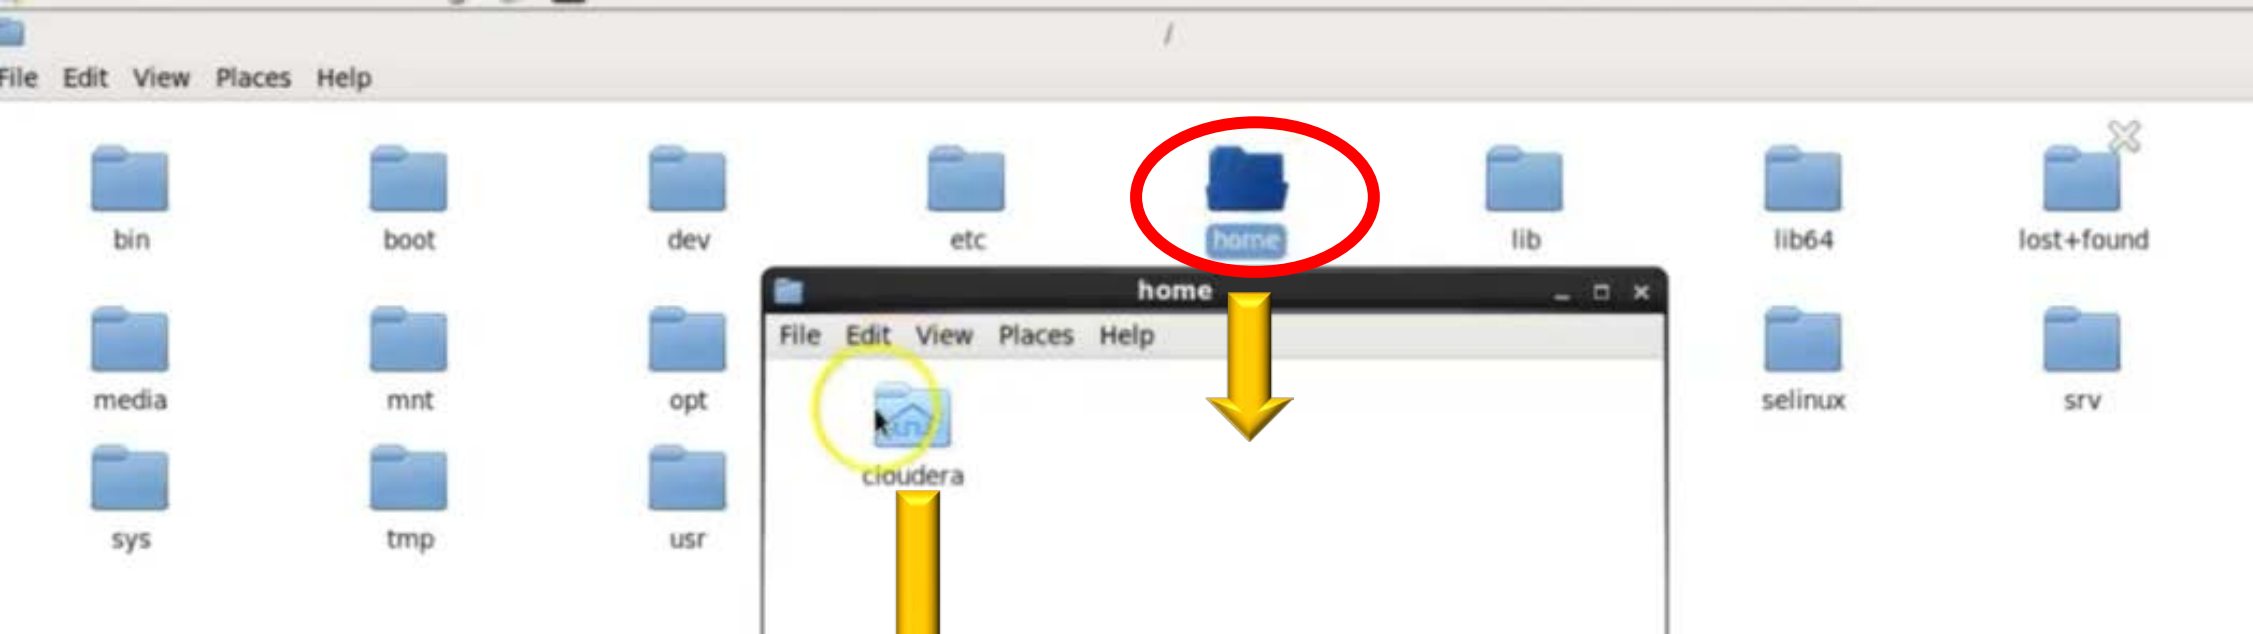
Computer














```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ ls  
cloudera-manager  eclipse  Music  testing  
cm_api.py         enterprise-deployment.json  parcels  Videos  
Desktop          express-deployment.json    Pictures  WordCount.jar  
Documents        kerberos  Public  workspace  
Downloads        lib       Templates  
[cloudera@quickstart ~]$ pwd  
/home/cloudera  
[cloudera@quickstart ~]$ cat > /home/cloudera/Processfile1.txt  
Riyaz Ahamed  
Riyaz Ahamed  
Riyaz Ahamed  
Cloudera  
Cloudera  
Cloudera  
Big Data  
Big Data  
Big Data  
Cloudera  
Riyaz Ahamed  
Riyaz Ahamed  
^Z  
[1]+  Stopped                  cat > /home/cloudera/Processfile1.txt  
[cloudera@quickstart ~]$
```

```
[cloudera@quickstart ~]$ hdfs dfs -ls /
Found 10 items
drwxrwxrwx - hdfs supergroup 0 2017-04-05 04:27 /benchmarks
drwxr-xr-x - hbase supergroup 0 2020-05-04 16:52 /hbase
drwxr-xr-x - cloudera supergroup 0 2020-05-01 02:57 /inputfolder
drwxr-xr-x - cloudera supergroup 0 2020-05-01 00:09 /inputnew
drwxr-xr-x - cloudera supergroup 0 2020-05-01 03:03 /output1
drwxr-xr-x - cloudera supergroup 0 2020-05-01 00:19 /output_new
drwxr-xr-x - solr solr 0 2017-04-05 04:29 /solr
drwxrwxrwt - hdfs supergroup 0 2020-04-23 22:28 /tmp
drwxr-xr-x - hdfs supergroup 0 2017-04-05 04:29 /user
drwxr-xr-x - hdfs supergroup 0 2017-04-05 04:29 /var
[cloudera@quickstart ~]$
```

```
[cloudera@quickstart ~]$ hdfs dfs -mkdir /inputfolder1
[cloudera@quickstart ~]$ hdfs dfs -put /home/cloudera/Processfile1.txt /inputfolder1/
[cloudera@quickstart ~]$ hdfs dfs -cat /inputfolder1/Processfile1.txt
Riyaz Ahamed
Riyaz Ahamed
Riyaz Ahamed
Cloudera
Cloudera
Cloudera
Big Data
Big Data
Big Data
Cloudera
Riyaz Ahamed
Riyaz Ahamed
[cloudera@quickstart ~]$
```

File Edit View Search Terminal Help

```
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/WordCount.jar WordCount /inputfolder1/Processfile1.txt /out1
20/05/04 17:41:39 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
20/05/04 17:41:43 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your ap
plication with ToolRunner to remedy this.
20/05/04 17:41:44 INFO input.FileInputFormat: Total input paths to process : 1
20/05/04 17:41:45 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:951)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:689)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:878)
20/05/04 17:41:45 INFO mapreduce.JobSubmitter: number of splits:1
20/05/04 17:41:46 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1588635643957_0003
20/05/04 17:41:48 INFO impl.YarnClientImpl: Submitted application application_1588635643957_0003
20/05/04 17:41:49 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1588635643957_0003/
20/05/04 17:41:49 INFO mapreduce.Job: Running job: job_1588635643957_0003
20/05/04 17:42:32 INFO mapreduce.Job: Job job_1588635643957_0003 running in uber mode : false
20/05/04 17:42:32 INFO mapreduce.Job:  map 0% reduce 0%
20/05/04 17:43:12 INFO mapreduce.Job:  map 100% reduce 0%
20/05/04 17:43:38 INFO mapreduce.Job:  map 100% reduce 100%
20/05/04 17:43:40 INFO mapreduce.Job: Job job_1588635643957_0003 completed successfully
```



```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
20/05/04 17:43:41 INFO mapreduce.Job: Counters: 49  
  File System Counters  
    FILE: Number of bytes read=67  
    FILE: Number of bytes written=242435  
    FILE: Number of read operations=0  
    FILE: Number of large read operations=0  
    FILE: Number of write operations=0  
    HDFS: Number of bytes read=254  
    HDFS: Number of bytes written=41  
    HDFS: Number of read operations=6  
    HDFS: Number of large read operations=0  
    HDFS: Number of write operations=2  
  Job Counters  
    Launched map tasks=1  
    Launched reduce tasks=1  
    Data-local map tasks=1  
    Total time spent by all maps in occupied slots (ms)=36101  
    Total time spent by all reduces in occupied slots (ms)=23876  
    Total time spent by all map tasks (ms)=36101  
    Total time spent by all reduce tasks (ms)=23876  
    Total vcore-seconds taken by all map tasks=36101  
    Total vcore-seconds taken by all reduce tasks=23876  
    Total megabyte-seconds taken by all map tasks=36967424  
    Total megabyte-seconds taken by all reduce tasks=24449024  
  Map-Reduce Framework  
    Map input records=12  
    Map output records=20  
    Map output bytes=208  
    Map output materialized bytes=67  
    Input split bytes=126  
    Combine input records=20  
    Combine output records=5  
    Reduce input groups=5  
    Reduce shuffle bytes=67  
    Reduce input records=5  
    Reduce output records=5  
    Spilled Records=10  
    Shuffled Maps =1  
    Failed Shuffles=0  
    Merged Map outputs=1  
    GC time elapsed (ms)=945  
    CPU time spent (ms)=4550  
    Physical memory (bytes) snapshot=337055744  
    Virtual memory (bytes) snapshot=3008081920  
    Total committed heap usage (bytes)=226365440  
  Shuffle Errors  
    BAD_ID=0  
    CONNECTION=0
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
Total vcore-seconds taken by all reduce tasks=23876  
Total megabyte-seconds taken by all map tasks=36967424  
Total megabyte-seconds taken by all reduce tasks=24449024  
Map-Reduce Framework  
Map input records=12  
Map output records=20  
Map output bytes=208  
Map output materialized bytes=67  
Input split bytes=126  
Combine input records=20  
Combine output records=5  
Reduce input groups=5  
Reduce shuffle bytes=67  
Reduce input records=5  
Reduce output records=5  
Spilled Records=10  
Shuffled Maps =1  
Failed Shuffles=0  
Merged Map outputs=1  
GC time elapsed (ms)=945  
CPU time spent (ms)=4550  
Physical memory (bytes) snapshot=337055744  
Virtual memory (bytes) snapshot=3008081920  
Total committed heap usage (bytes)=226365440  
Shuffle Errors  
BAD_ID=0  
CONNECTION=0
```

```
[cloudera@quickstart ~]$ hdfs dfs -ls /out1
Found 2 items
-rw-r--r-- 1 cloudera supergroup          0 2020-05-04 17:43 /out1/_SUCCESS
-rw-r--r-- 1 cloudera supergroup    41 2020-05-04 17:43 /out1/part-r-00000
[cloudera@quickstart ~]$
```

```
[cloudera@quickstart ~]$ hdfs dfs -cat /out1/part-r-00000
Ahamed 5
Big 3
Cloudera 4
Data 3
Riyaz 5
[cloudera@quickstart ~]$
```

```
[cloudera@quickstart ~]$ cat /home/cloudera/Processfile1.txt
Riyaz Ahamed
Riyaz Ahamed
Riyaz Ahamed
Cloudera
Cloudera
Cloudera
Big Data
Big Data
Big Data
Cloudera
Riyaz Ahamed
Riyaz Ahamed
[cloudera@quickstart ~]$
```



Référence

<https://hadoop.apache.org/docs/stable/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>